

Target–Deployment Routed Refinement: Broadening the Stability Envelope of Offline TD3+BC

Anonymous submission

Abstract

TD3+BC ordinarily uses one actor branch both to construct bootstrapped critic targets and to produce the deployed policy, tying target-policy displacement to the total reach of policy improvement. We introduce target–deployment routed refinement (TDRR), a staged TD3+BC actor chain whose first-hop target copy supplies TD actions and whose final hop is deployed. The squared action-anchored objective has a deterministic Wasserstein–2 proximal interpretation. We derive an inexact finite-step margin relative to the learned critic, an explicit critic-error term, a first-hop target-perturbation bound, and a fixed- K local expansion. Across nine D4RL locomotion tasks and fourteen fixed budgets, three proximal hops reduce high-budget low-return cells from 35/63 to 17/63; a scale-matched linearized realization reduces them to 24/63. Three-hop TDRR lowers sample-anchored next-state target-policy displacement in 88/90 paired runs while lowering the current-state final-policy displacement proxy in only 36/90. These results support an empirical stability-envelope gain for the implemented refinement bundle without attributing it to routing alone.

Introduction

Offline reinforcement learning (RL) seeks a policy from a fixed dataset without additional environment interaction. Its central difficulty is not merely to select a high-value action, but to determine how far policy improvement can proceed before the learned critic becomes unreliable. TD3+BC regularizes a deterministic actor toward dataset actions while using a Polyak copy of that actor in bootstrapped critic targets (Fujimoto and Gu 2021). The same policy branch therefore determines both the action shift entering critic training and the reach of the evaluated policy. Aggressive improvement can move target actions into regions where critic error is amplified, whereas conservative improvement can leave useful policy improvement unrealized.

We study this coupling through target–deployment routed refinement (TDRR), a re-centered actor architecture for TD3+BC. TDRR represents a nominal improvement budget with K sequential actor updates. The first actor is anchored to dataset actions, and each later actor is anchored to the stop-gradient output of

its predecessor. A target copy of the first refined actor supplies the next actions in TD targets, while the final actor is used for evaluation. The architecture thus assigns critic-target construction and deployment to distinct policy branches without changing the critic update rule.

Let D_{critic} denote the squared displacement of the critic-coupled first-hop branch and let D_{final} denote that of the evaluated final policy. A one-hop actor identifies these quantities; TDRR assigns them to μ_1 and μ_K , respectively. This separation can increase final-policy reach at comparable first-hop movement or contract both movements near an unstable training outcome. It also introduces a countervailing mismatch because the critic target follows the μ_1 branch while evaluation follows μ_K . The method therefore exposes a design axis between target-policy displacement and deployment reach rather than guaranteeing that either quantity is independently optimal.

For deterministic actors, the squared action anchor in TD3+BC is a Wasserstein–2 movement cost in the Dirac-policy limit. Under a fixed critic, inactive action projection, accurate subproblem solutions, and small realized steps, explicit and implicit updates are first-order discretizations of the same Wasserstein gradient flow. At finite steps, a proximal comparison yields an improvement margin only relative to the learned critic. At large nominal budgets, projection, neural actor fitting, critic drift, saturation, and the seeded optimization trajectory determine the observed end-to-end outcome. We keep these local, finite-step, and coupled regimes distinct throughout the analysis.

We evaluate this formulation with a complete fixed-budget sweep over nine D4RL locomotion tasks, fourteen budgets, and two paired seeds. Three proximal hops reduce large-budget low-return cells from 35/63 for TD3+BC to 17/63 and primarily improve the lower tail rather than uniformly shifting scores. Three scale-matched projected linearized hops reduce failures to 24/63, so the stability effect is not unique to the proximal loss. Direct next-state diagnostics show that three-hop TDRR lowers deterministic target-policy displacement in 88/90 paired comparisons with TD3+BC, while its current-state final-policy displacement proxy is lower

in only 36/90. Low-return runs additionally exhibit orders-of-magnitude increases in TD-error tails, larger realized movements, and critic gains that reverse sign under a controlled common critic. Targeted route and compute controls probe components of the architecture but do not identify the source of the sweep-wide gain.

Our contributions are as follows.

- We formulate TDRR for TD3+BC as a re-centered actor composition that assigns target-policy displacement and final-policy reach to distinct branches.
- We give a Wasserstein proximal interpretation, derive an inexact finite-step margin with an explicit critic-error term, bound first-hop TD-target perturbation, and separate these statements from the fixed-critic small-step flow limit.
- Through complete proximal and scale-matched linearized sweeps, direct target-policy measurements, critic-failure diagnostics, and targeted controls, we show that the tested three-hop refinement bundles broaden the empirical large-budget stability envelope and that low-return outcomes are consistent with critic miscalibration in the nonlocal regime.

Related Work

Behavior-regularized offline actor–critic learning. Offline actor–critic methods must control value extrapolation when a learned policy selects actions outside the fixed dataset. Prior work studies this failure through constrained policy improvement, conservative value learning, and restrictions on actions used in Bellman backups (Fujimoto, Meger, and Precup 2019; Kumar et al. 2019). TD3+BC provides a particularly minimal control: a deterministic actor maximizes a normalized critic while a squared cloning term anchors its actions to the dataset (Fujimoto and Gu 2021). We keep this critic update and actor family fixed and study which refined actor branch supplies the TD target.

Policy geometry and proximal refinement. Trust-region policy optimization is commonly derived from information geometry (Amari and Nagaoka 2000; Schulman et al. 2015). Optimal-transport geometry gives a different notion of movement: policy optimization can be expressed as a Wasserstein gradient flow, with related natural-gradient and trust-region constructions for continuous control (Zhang et al. 2018; Moskovitz et al. 2020; Terpin et al. 2022; Song et al. 2023). For deterministic policies, the statewise Wasserstein–2 cost reduces to squared action displacement. We use that identity to interpret the TD3+BC anchor as a proximal movement cost, while separating the finite-step variational statement from its small-step flow limit.

Target networks and actor–critic coupling. Bootstrapping error can feed critic extrapolation back into actor updates (Fujimoto, Meger, and Precup 2019; Kumar et al. 2019). Polyak averaging delays this feedback but

ordinarily copies the same actor branch that will be deployed. Target–deployment routed refinement makes a narrower architectural change: the target copy follows the first local refinement, whereas evaluation uses the final refinement. Our target-perturbation bound and direct displacement audit therefore concern branch assignment inside TD3+BC rather than a new critic objective or a general claim about offline policy extraction.

Preliminaries

We consider a discounted Markov decision process $\langle \mathcal{S}, \mathcal{A}, P, r, \gamma \rangle$ with continuous actions and a fixed offline dataset $\mathcal{D} = \{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$. A clipped double critic is trained from targets

$$y_\rho = r + \gamma \min_j Q_j^-(s', \rho(s')), \quad (1)$$

where ρ is a target policy. When $\rho(s')$ leaves dataset support, critic error enters the regression target and can be amplified by later actor updates (Fujimoto, Meger, and Precup 2019; Kumar et al. 2019).

For a deterministic actor μ_θ , normalized TD3+BC minimizes

$$\mathcal{L}_{\text{actor}}(\theta) = -\alpha \mathbb{E}_s \bar{Q}(s, \mu_\theta(s)) + \mathbb{E}_{(s,a) \sim \mathcal{D}} \frac{\|\mu_\theta(s) - a\|_2^2}{d}, \quad (2)$$

where $\bar{Q} = \hat{Q}/C$ and $C > 0$ is a stop-gradient critic scale (Fujimoto and Gu 2021). Ordinarily, a Polyak copy of this actor enters Eq. (1), while the online actor is deployed. Target-policy movement and deployment reach therefore follow the same branch.

A metric proximal update with step $h > 0$ takes the form

$$\pi^+ \in \arg \min_{\pi} \left\{ \hat{\mathcal{E}}(\pi) + \frac{1}{2h} d_{\mathcal{M}}^2(\pi, \pi^{\text{ref}}) \right\}. \quad (3)$$

This minimizing-movement construction is associated with metric gradient flows (Ambrosio, Gigli, and Savaré 2005; Otto 2001). For deterministic state-conditioned policies,

$$W_2^2(\delta_{\mu(s)}, \delta_{\nu(s)}) = \|\mu(s) - \nu(s)\|_2^2. \quad (4)$$

Thus the TD3+BC anchor is a Wasserstein–2 movement cost up to fixed scaling. We distinguish its finite-step variational meaning from the small-step, fixed-critic gradient-flow limit and from large-budget coupled training.

Target–Deployment Routed Refinement

At each delayed actor update, TDRR maintains K independently initialized, persistent deterministic actors with separate optimizer states. The first actor takes one TD3+BC update toward the sampled dataset actions; each later actor takes one update toward the stop-gradient output of its predecessor on the same minibatch. The online critic is fixed within this chain; we denote its first Q-function by $\hat{Q}$. A Polyak copy of the

first actor supplies TD-target actions, and the final actor is evaluated. We first derive the proximal operator underlying these updates and then formalize their composition.

Offline actor updates as proximal improvement

Let $\rho_{\mathcal{D}}$ be the empirical state distribution and consider a deterministic actor $\mu : \mathcal{S} \rightarrow \mathcal{A} \subset \mathbb{R}^d$. At one training iteration, the critic parameters are held fixed while the actor chain is updated. We write the current critic as $\hat{Q}$ and normalize its global scale by

$$\bar{Q}_k(s, a) = \frac{\hat{Q}(s, a)}{C_k}, \quad (5)$$

$$C_k = \text{sg} \left(\mathbb{E}_s |\hat{Q}(s, \tilde{\mu}_k(s))| + \epsilon \right).$$

where $\tilde{\mu}_k$ is the actor before hop k and sg denotes stop-gradient. A fixed positive C only rescales gradient-flow time; recomputing C_k changes the finite-step schedule. Our finite-step statements condition on the realized scales.

For deterministic policies and a state distribution ρ , define the action-averaged Wasserstein metric

$$\mathbf{d}_{d,\rho}^2(\mu, \nu) := \mathbb{E}_{s \sim \rho} \left[\frac{1}{d} \|\mu(s) - \nu(s)\|_2^2 \right]. \quad (6)$$

It is d^{-1} times the state-averaged squared W_2 distance between state-conditioned Dirac policies. We abbreviate $\mathbf{d}_d = \mathbf{d}_{d,\rho_{\mathcal{D}}}$ on the actor-training state distribution. The population counterpart of squared behavior cloning also admits a deterministic reference. Let P_B be the data-generating behavior distribution, let ρ_B be its state marginal, and define $b(s) = \mathbb{E}_{P_B}[a | s]$. Then

$$\mathbb{E}_{(s,a) \sim P_B} \frac{1}{d} \|\mu(s) - a\|_2^2 = \mathbf{d}_{d,\rho_B}^2(\mu, b) + \text{const}, \quad (7)$$

where the constant is the conditional action variance and does not depend on μ . The empirical actor loss is a Monte Carlo estimate of the left-hand side, so samplewise squared anchoring targets the conditional-mean behavior map at the population level.

Given a reference policy ν and step $h > 0$, define the proximal refinement operator by

$$\mathcal{P}_h^{\hat{Q}}(\nu) \in \arg \min_{\mu \in \mathcal{M}} \left\{ -\mathbb{E}_s \bar{Q}_k(s, \mu(s)) + \frac{1}{2h} \mathbf{d}_d^2(\mu, \nu) \right\}. \quad (8)$$

Multiplying by $2h$ gives the practical actor objective

$$\mathcal{L}_{h,k}(\mu | \nu) = -2h \mathbb{E}_s \bar{Q}_k(s, \mu(s)) + \mathbf{d}_d^2(\mu, \nu). \quad (9)$$

Proposition 1 (Behavior anchoring is Wasserstein proximal improvement). For deterministic policies, the normalized anchored objective

$$-\alpha \mathbb{E}_s \bar{Q}(s, \mu(s)) + \mathbf{d}_d^2(\mu, \nu)$$

is Eq. (8), up to multiplication by a positive constant, with $h = \alpha/2$. Equivalently, the raw-critic objective $-\lambda \mathbb{E}_s \hat{Q}(s, \mu(s)) + \mathbf{d}_d^2(\mu, \nu)$ has $h = \lambda C/2$ when C is fixed. The population counterpart of squared dataset anchoring uses $\nu = b$ by Eq. (7).




Figure 1: TDRR re-centers each actor update. A Polyak copy of μ_1 supplies TD-target actions, while later actors increase or reshape final-policy reach without entering the target policy.

The proposition identifies an update operator relative to the estimated critic; true-return improvement lies outside the scope of this certificate.

Re-centered composition and actor-critic dataflow

For population analysis, let $\mu_0 := b$ denote the conditional-mean behavior map and let $h_{1:K}$ be positive step sizes. The idealized TDRR chain is

$$\mu_k \in \mathcal{P}_{h_k}^{\hat{Q}}(\text{sg } \mu_{k-1}), \quad k = 1, \dots, K. \quad (10)$$

The implementation approximates each subproblem with one optimizer update. Its first actor μ_1 uses sampled dataset actions as its reference; later actors $\mu_2, \dots, \mu_K$ use the preceding online actor under the same current critic. Only the Polyak copy of μ_1 enters the TD target, while μ_K supplies the evaluated action.

Our fixed-budget schedule sets $h_k = T/K$ and holds the total nominal budget T fixed. Increasing K jointly changes local step size, re-centering frequency, actor compute, and the target/deployment dataflow; the experiment therefore estimates the effect of this implemented refinement bundle.

On the next-state distribution $\rho'_{\mathcal{D}}$ used by the TD target, define the action-averaged population displacements

$$D_{\text{critic}} := \mathbf{d}_{d,\rho'_{\mathcal{D}}}^2(\mu_1, \mu_0), \quad (11)$$

$$D_{\text{final}} := \mathbf{d}_{d,\rho'_{\mathcal{D}}}^2(\mu_K, \mu_0).$$

A one-step actor has $D_{\text{critic}} = D_{\text{final}}$; TDRR assigns the two population quantities to distinct policy branches. On fixed validation batches, we measure the sample-anchored online branches and the deterministic next-state Polyak target branch:

$$\hat{D}_{\text{critic}} = \mathbb{E}_i \|\mu_1(s_i) - a_i\|_2^2/d,$$

$$\hat{D}_{\text{final}} = \mathbb{E}_i \|\mu_K(s_i) - a_i\|_2^2/d, \quad (12)$$

$$\hat{D}_{\text{target}} = \mathbb{E}_i \|\mu_1^-(s'_i) - a_{i+1}\|_2^2/d.$$

The direct audit retains only bootstrapped transition pairs; a smoothed version uses common TD3 target-noise draws. Sample anchoring adds a common conditional-action-variance floor in population differences. These are movement measurements, not critic-reliability estimates.

Finite-step statements

The following statements do not require a small-step Taylor expansion.

Proposition 2 (Inexact estimated-energy margin). Suppose μ_{k-1} is feasible for the k th actor class and the computed actor satisfies

$$\mathcal{L}_{h_k,k}(\mu_k \mid \mu_{k-1}) \leq \mathcal{L}_{h_k,k}(\mu_{k-1} \mid \mu_{k-1}) + \varepsilon_k.$$

Then

$$m_k^{\hat{Q}} := \frac{2h_k}{C_k} \mathbb{E}_s[\hat{Q}(s, \mu_k(s)) - \hat{Q}(s, \mu_{k-1}(s))] - d_d^2(\mu_k, \mu_{k-1}) \geq -\varepsilon_k. \quad (13)$$

For the first sample-anchored hop, a statewise dataset action is not generally a feasible deterministic-policy comparator; the population reference in Eq. (7) or a behavior actor is required for a literal certificate. For re-centered hops, the preceding actor is feasible by construction.

The margin is relative to $\hat{Q}$. Let a trusted objective be $Q^\dagger$ and write $\hat{Q} = Q^\dagger + e$. Then

$$m_k^{Q^\dagger} = m_k^{\hat{Q}} - \frac{2h_k}{C_k} \mathbb{E}_s[e(s, \mu_k(s)) - e(s, \mu_{k-1}(s))]. \quad (14)$$

Thus estimated-energy descent transfers to $Q^\dagger$ only when critic error is controlled along the realized path. Equation (13) certifies the learned objective rather than monotone true return.

Proposition 3 (First-hop TD-target perturbation bound). For clarity, consider a noiseless TD target using μ_1 , $y_{\mu_1} = r + \gamma \min_j Q_j^-(s', \mu_1(s'))$, and suppose each target critic $Q_j^-(s', \cdot)$ is $L(s')$ -Lipschitz. For any reference μ_0 ,

$$\|y_{\mu_1} - y_{\mu_0}\|_{L_2(\mathcal{D})} \leq \gamma \left(\mathbb{E}_{s'} L(s')^2 \cdot \|\mu_1(s') - \mu_0(s')\|_2^2 \right)^{1/2}. \quad (15)$$

If $L(s') \leq \bar{L}$, the right-hand side is at most $\gamma \bar{L} \sqrt{d D_{\text{critic}}}$. Conditional on μ_1 , later actors do not enter this one-step target perturbation.

Proposition 3 bounds a one-step value-target difference, not the stability of the learned critic. Its architectural content is that later refinement can alter D_{final} without directly entering this bound. With a Polyak target network, μ_1 is replaced by its target copy. A same-noise comparison under TD3 target smoothing obeys the same bound because action clipping is nonexpansive. The current-state $\hat{D}_{\text{critic}}$ omits Polyak lag and target noise, whereas $\hat{D}_{\text{target}}$ measures the Polyak actor directly and has a smoothed counterpart. Neither displacement measures the action Lipschitz factor or the resulting change in critic value.

The separation introduces a complementary evaluation-policy mismatch: critic targets follow

the μ_1 branch, whereas evaluation follows μ_K . Increasing D_{final} at fixed D_{critic} can therefore make the critic less representative of the evaluated policy. We interpret TDRR as exposing a measurable target/deployment trade-off; its empirical return captures the net effect of target movement, evaluation mismatch, actor compute, and the coupled optimization trajectory.

Local Wasserstein-flow limit

The gradient-flow interpretation is a local organizing principle. Fix a critic and an interior action a_0 . Let $g = \nabla_a \hat{Q}(a_0)$, $H = \nabla_a^2 \hat{Q}(a_0)$, and $\beta = dh/C$. Before projection, the scale-matched linearized and exact statewise proximal updates satisfy

$$\Delta a_{\text{lin}} = \beta g, \quad \Delta a_{\text{prox}} = \beta \nabla_a \hat{Q}(a_0 + \Delta a_{\text{prox}}). \quad (16)$$

When $\beta \|H\| \ll 1$, the latter yields

$$\Delta a_{\text{prox}} = (I - \beta H)^{-1} \beta g = \beta g + \beta^2 H g + O(\beta^3), \quad (17)$$

and both updates are first-order discretizations of $\dot{a} = d \nabla_a \hat{Q}(a)$. With fixed C this is a constant time-rescaling of raw-critic gradient ascent; with a positive policy-dependent global C it follows the same direction under an adaptive time parameterization. This expansion additionally requires small realized displacement, inactive projection, and a sufficiently accurate actor solve. At large budgets, our experiments diagnose the resulting nonlocal actor-critic regime instead of extrapolating Eq. (17). The end-to-end return sweep, even at small nominal T , is not a frozen-critic action-level consistency test of this limit.

Proposition 4 (Local fixed-budget composition). Let a smooth one-step update in local coordinates satisfy $\Psi_h(x) = x + hf(x) + h^2 p(x) + O(h^3)$. For fixed K and $T \rightarrow 0$,

$$\begin{aligned} \Psi_{T/K}^K(x_0) &= x_0 + T f_0 \\ &\quad + T^2 \left[\frac{K-1}{2K} J_0 f_0 + \frac{1}{K} p_0 \right] + O(T^3). \end{aligned} \quad (18)$$

where $f_0 = f(x_0)$ and $J_0 = Df(x_0)$. Coordinate explicit Euler has $p_0 = 0$, whereas exact backward Euler has $p_0 = J_0 f_0$. Hence both are first-order methods; for each fixed K , the $T \rightarrow 0$ expansion has a leading flow-bias coefficient proportional to $1/K$.

Proposition 4 explains a possible fixed-critic small-step benefit of substepping. It neither predicts the sign of true-return change nor applies to the projected, saturated, moving-critic branches studied at large T . The expansion is not a uniform $K \rightarrow \infty$ result.

Experiments

We evaluate TDRR on the nine D4RL locomotion tasks formed by HalfCheetah, Hopper, and Walker2d with medium, medium-replay, and expert datasets (Fu

et al. 2020). These tasks directly exercise deterministic continuous control and provide distinct behavior-data regimes for a target-routing stress test. Base implementations follow CORL configurations (Tarasov et al. 2022); scores are D4RL normalized returns. Additional per-task results and diagnostic definitions are in the Supplementary Material.

Fixed-budget protocol and estimand. We hold a total nominal budget T fixed, split it across K actor hops as $h = T/K$, and sweep

$$T \in \{0.05, 0.1, 0.2, 0.4, 0.7, 1.5, 2.5, 4, 7, 10, 12, 14, 17, 20\}.$$

over all nine tasks with two paired seeds. The sweep is complete over this specified grid and contains $9 \times 14 \times 2$ runs per method. It is designed as a high-resolution stability phase diagram rather than a final benchmark alone. For $K = 1$, Proposition 1 maps T to the normalized TD3+BC coefficient $\alpha = 2T$.

We compare TD3+BC ($K = 1$), proximal TDRR with $K \in \{2, 3\}$ (TDRR-Prox), and a scale-matched projected linearized target with $K \in \{2, 3\}$ (TDRR-Lin):

$$a_{k+1}^{\text{tar}} = \Pi_{\mathcal{A}} \left[a_k + d \frac{T}{K} \frac{\nabla_a \hat{Q}(s, a_k)}{C_k} \right]. \quad (19)$$

Here a_0 is the sampled dataset action. At each hop, TDRR-Lin detaches the current reference, recomputes the $\hat{Q}$ gradient and C_k , clips the target to the action box, and takes one mean-squared-error optimizer step with the corresponding persistent actor. It uses the same first-hop target and final-hop evaluation routing as TDRR-Prox. The factor d matches the local action-space coefficient induced by the action-averaged metric in Eq. (6). All conditions match datasets, critic updates, and the horizon of 10^6 critic updates; checkpoint diagnostics use the final checkpoint. A K -hop condition performs K actor optimizer steps at each delayed policy update, and C_k is recomputed at each reference. The primary estimand is the end-to-end effect of the fixed-budget refinement bundle, which jointly changes local step size, re-centering frequency, actor compute, and target/deployment dataflow.

The analyses test whether this bundle reduces large-budget low-return outcomes, whether the effect persists under the linearized realization, how first-hop target-policy displacement differs from final-policy reach, and which critic-validity diagnostics characterize failed training trajectories.

Three-hop refinement broadens the stability envelope

Figure 2 and Table 1 summarize the fixed-budget sweep. We descriptively call a high-budget cell collapsed when its two-seed mean score is below 20; this is a locomotion failure indicator, not a theoretical threshold. TDRR-Prox with $K = 3$ has the strongest high-budget average

Method	$T \leq 2.5$	High	Tail	Collapse
TD3+BC	68.35	25.97	20.46	35 / 68
TDRR-P2	71.62	42.02	31.49	25 / 53
TDRR-P3	70.56	52.78	43.55	17 / 41
TDRR-L2	69.02	36.68	27.44	29 / 58
TDRR-L3	68.54	43.91	33.12	24 / 53

Table 1: Aggregate fixed-budget results. High denotes $T \geq 4$ and tail denotes $T \geq 10$. Collapse entries are seed-mean cells / raw seeded runs; P and L denote proximal and linearized.

and reduces the seed-mean collapse count from 35/63 to 17/63. TDRR-Lin with $K = 3$ reduces the same count to 24/63, compared with 29/63 for two linearized hops, and raises the high-budget average from 36.68 to 43.91. The raw-seed counts confirm the lower failure rate of both three-hop realizations. The proximal $K = 2$ to $K = 3$ mean improvement for $T \geq 4$ is +10.76, whereas the median is only +0.50; the corresponding linearized differences are +7.23 and +0.86. Thus additional hops primarily rescue low-return outcomes rather than uniformly shifting settings. The Supplementary Material reports sensitivity to thresholds 0, 10, 20, 30, and 40.

Task-dependent response and a fourth hop. The budget response is smooth on HalfCheetah-medium, plateau-like on Hopper-medium-replay, sharply non-monotone on Hopper-medium and Walker2d-medium, and asymmetric on expert datasets. Within-task transitions are often organized more clearly by realized first-hop movement than by nominal T , but comparable movement can yield very different returns across tasks. There is therefore no universal critic-valid radius. In a separate four-seed Hopper-medium study, TDRR-Prox with $K = 4$ scores 100.9 ± 1.7 at $T = 17$, where $K = 3$ scores 0.7, but becomes highly variable at $T = 20$ (41.8 ± 44.0). The fourth hop relocates a nonmonotone high-return window; it does not establish a linear frontier law or identify a local-step mechanism.

The linearized sweep shows that the gain is not unique to the proximal loss. At high budgets, TDRR-Lin with $K = 3$ gains 7.23 points over $K = 2$ and reduces collapse cells from 29 to 24. TDRR-Prox remains stronger overall, so these are comparisons of complete training procedures rather than isolated integrator effects.

Large-budget behavior is nonlocal and critic dependent

At $T = 20$, the raw linearized target reaches the action box in 94–95% of states, and the effective action-step coefficient h_{eff} is only 14% and 9% of the nominal dT/K coefficient in the two hops. These are medians over 18 environment–seed runs per hop. The practical operation is projected target generation followed by amortized neural regression under a moving critic. The sweep




Figure 2: Fixed-budget sweep over nine locomotion tasks. (a) Mean score over tasks. (b) Collapse rates for $T \geq 4$, reported both by environment-budget cell and by seeded run. The tested refinement bundles broaden the stable high-budget region for both proximal and scale-matched linearized realizations.

measures the resulting end-to-end stability envelope; local truncation error lies outside this estimand.

Figure 3 shows the common failure signature. Environment return supplies the external label: collapsed runs score below 20. These runs have orders-of-magnitude larger TD-error tails, larger critic magnitudes, and larger realized movements. The controlled common critic is the frozen final TD3+BC $T = 1$ critic matched by environment and seed. A run-total gain sums the mean Q change across actor hops on the shared state batch. For collapsed TDRR-Lin runs, its median under the jointly trained critic is 3.33×10^{11} ; under the common critic it is -1.42 , compared with $+2.11$ for stable runs. For TDRR-Prox, the common-critic gain changes from $+2.51$ in stable runs to -1.06 in collapsed runs. This sign reversal diagnoses cross-critic disagreement consistent with Eq. (14); environment return remains the criterion for policy failure.

A simulator-reference CPU audit provides targeted evidence of critic miscalibration in the screened checkpoints: the median signed error $\hat{Q}_1(s, a_K) - Q_{MC}^{P_1}(s, a_K)$ is $+1.70 \times 10^{12}$ in five collapsed checkpoints and -38.7 in 55 stable ones. Its action-ranking subset is sparse, so we use it only as targeted checkpoint-level evidence.

Target-policy displacement and final-policy reach

At the same nominal budget, TDRR-Prox realizes a smaller current-state first-hop displacement proxy than TDRR-Lin once $T \geq 7$ (the Supplementary Material). At $T = 20$, their median $\hat{D}_{\text{critic}}$ values are 0.195 and 0.323. Across 90 paired seeded comparisons at $T \in \{4, 7, 10, 14, 20\}$, exactly one realization collapses in 11 cases. In all 11, the stable realization has both smaller $\hat{D}_{\text{critic}}$ and smaller $\hat{D}_{\text{final}}$. These paired outcomes associate smaller realized movement with the stable optimization trajectory. Because the proxy is measured on current validation states rather than the Polyak tar-

get policy on next states, this result organizes the observed failure transition without identifying a causal route through the TD target.

We directly measure Eq. (12) on all 270 audited cells. The current-state proxy closely tracks deterministic target-policy displacement ($r = .9999$). TDRR-Prox $K = 3$ has lower target-policy displacement than TD3+BC in 88 of 90 paired runs (median ratio .658), while the current-state $\hat{D}_{\text{final}}$ proxy is lower in only 36 of 90 (median ratio 1.057). Thus the three-hop refinement can reduce movement in the target-policy branch without merely shrinking its measured final-policy reach. Against $K = 2$, $K = 3$ lowers target-policy displacement in 84 of 90 runs but loses by more than one score point in 17 of those 84, showing that target-policy displacement alone does not determine performance (the Supplementary Material).

A targeted $K = 3$ study on Hopper-medium reveals two modes of decoupling (the Supplementary Material). For $T = 4-10$, $K = 2$ and $K = 3$ have nearly identical $\hat{D}_{\text{critic}}$, whereas $K = 3$ increases $\hat{D}_{\text{final}}$ and score; at $T = 10$, the scores are 58.12 and 85.35. In the observed transition region from stable to collapsed scores ($T = 12, 14$), $K = 3$ instead contracts both displacements and avoids the low-return outcome. At $T = 20$, both enter severe projection and saturation and collapse. The third hop changes target-policy movement, re-centering frequency, and the optimization trajectory rather than the formal order of the local update.

Targeted component controls

We next hold the actor trajectory fixed and change only the backup route. Two actor-detached shadow critics start from the same parameters and receive identical minibatches and target-smoothing noise; one bootstraps through the first-hop target actor and the other through the final-hop target actor. The study covers




Figure 3: Median diagnostics over the matched $T \in \{4, 7, 10, 14, 20\}$ audit. Proximal and linearized realizations share a diagnostic pattern consistent with critic miscalibration: critic errors and realized movements increase sharply, and run-total gains under a common critic become negative. Environment return labels failure; the common critic diagnoses cross-critic disagreement.

Hopper-medium at $T \in \{10, 20\}$ and Walker2d-medium at $T \in \{4, 7\}$, with two seeds and four checkpoints. Relative to a common simulator return under the effective first-route backup continuation, the final-route shadow has greater state-median symmetric relative error in all eight run-level averages. The final-minus-first route delta has mean .124, median .013, a run-bootstrap mean 95% interval [.008, .323], and an exact two-sided sign-flip $p = .0078$. Because the final-route critic targets a different continuation, this contrast measures route-induced alignment with the primary first-route value estimand, not each shadow critic’s calibration to its own continuation. Neither shadow critic drives the actors, so the intervention does not estimate a routing effect on return.

Two additional eight-run controls probe simpler explanations (the Supplementary Material). With three actor optimizer calls in every condition, sequential $K = 3$, a fixed-reference chain, and a direct-final control have means 71.4, 70.0, and 59.4 and collapse counts 2/8, 2/8, and 3/8. Fixed-reference refinement therefore matches sequential re-centering in this targeted sample, so the experiment does not establish re-centering alone as the source of the sweep-wide gain. Slowing only the $K = 1$ actor-target Polyak rate by a factor of three changes the matched collapse count from 6/8 to 7/8 and does not reproduce the $K = 3$ rescue; actual displacement under this lag intervention was not audited. The route-only comparison characterizes alignment with the first-route estimand; together with the compute and lag controls, it refines the architectural diagnosis but does not identify routing as the cause of the sweep-wide gain.

Discussion and limitations

Both three-hop realizations reduce large-budget failures, although TDRR-Prox remains stronger overall. Direct measurements show that lower first-hop target movement need not shrink final-policy reach, and fixed-trajectory final-route shadows are less aligned with the

common first-route value estimand in all eight targeted runs. Fixed-reference refinement matches sequential re-centering, while slower target averaging does not rescue the selected one-hop cells. Thus the primary estimate remains the implemented bundle: K changes local step, re-centering, compute, normalization, and routing together. The fourth-hop study likewise relocates rather than uniformly improves the high-return window.

The complete grid covers TD3+BC on nine locomotion datasets with two paired seeds per cell. The separate-machine four-seed $K = 4$ study is represented by transferred score CSVs but lacks matched checkpoints and configuration manifests. Collapse is descriptive, and the common-critic, simulator, and eight-run controls are targeted rather than exhaustive. We do not provide an offline rule for selecting T or K . The finite-step margin concerns the learned critic; the Wasserstein-flow account assumes small realized steps under a fixed critic.

Conclusion

TDRR separates critic-target and deployment roles across a re-centered TD3+BC actor chain, with a deterministic Wasserstein proximal interpretation in the fixed-critic local regime. In the complete sweep, three proximal and linearized hops reduce large-budget low-return cells from 35/63 to 17/63 and 24/63, respectively. Three-hop TDRR lowers sample-anchored next-state target-policy displacement in 88/90 pairs without generally shrinking the current-state final-policy proxy. These results support the end-to-end claim that the tested three-hop refinement bundle broadens empirical stability while separating target exposure from deployment reach; they do not attribute the sweep-wide gain to routing, re-centering, or compute alone.

References

Amari, S.-i.; and Nagaoka, H. 2000. *Methods of Information Geometry*. American Mathematical Society.

Ambrosio, L.; Gigli, N.; and Savaré, G. 2005. Gradient Flows in Metric Spaces and in the Space of Probability Measures. Birkhäuser.

Fu, J.; Kumar, A.; Nachum, O.; Tucker, G.; and Levine, S. 2020. D4RL: Datasets for Deep Data-Driven Reinforcement Learning. arXiv:2004.07219.

Fujimoto, S.; and Gu, S. S. 2021. A Minimalist Approach to Offline Reinforcement Learning. In Advances in Neural Information Processing Systems.

Fujimoto, S.; Meger, D.; and Precup, D. 2019. Off-Policy Deep Reinforcement Learning without Exploration. In Proceedings of the 36th International Conference on Machine Learning.

Kumar, A.; Fu, J.; Tucker, G.; and Levine, S. 2019. Stabilizing Off-Policy Q-Learning via Bootstrapping Error Reduction. In Advances in Neural Information Processing Systems.

Moskowitz, T.; Arbel, M.; Huszár, F.; and Gretton, A. 2020. Efficient Wasserstein Natural Gradients for Reinforcement Learning. arXiv preprint arXiv:2010.05380.

Otto, F. 2001. The Geometry of Dissipative Evolution Equations: The Porous Medium Equation. Communications in Partial Differential Equations, 26(1–2): 101–174.

Schulman, J.; Levine, S.; Moritz, P.; Jordan, M. I.; and Abbeel, P. 2015. Trust Region Policy Optimization. In Proceedings of the 32nd International Conference on Machine Learning.

Song, J.; He, N.; Ding, L.; and Zhao, C. 2023. Provably Convergent Policy Optimization via Metric-Aware Trust Region Methods. arXiv preprint arXiv:2306.14133.

Tarasov, D.; Nikulin, A.; Akimov, D.; Kurenkov, V.; and Kolesnikov, S. 2022. CORL: Research-Oriented Deep Offline Reinforcement Learning Library. NeurIPS Offline Reinforcement Learning Workshop.

Terpin, A.; Lanzetti, N.; Yardim, B.; Dörfler, F.; and Ramponi, G. 2022. Trust Region Policy Optimization with Optimal Transport Discrepancies. arXiv preprint arXiv:2210.11137.

Zhang, R.; Chen, C.; Li, C.; and Carin, L. 2018. Policy Optimization as Wasserstein Gradient Flows. In Proceedings of the 35th International Conference on Machine Learning.

Reproducibility Checklist

Instructions for Authors:

This document outlines key aspects for assessing reproducibility. Please provide your input by editing this .tex file directly.

For each question (that applies), replace the “Type your response here” text with your answer.

Example: If a question appears as

```
\question{Proofs of all novel claims are
included} {(yes/partial/no)}
Type your response here
```

you would change it to:

```
\question{Proofs of all novel claims are
included} {(yes/partial/no)}
yes
```

Please make sure to:

- Replace ONLY the “Type your response here” text and nothing else.
- Use one of the options listed for that question (e.g., yes, no, partial, or NA).
- Not modify any other part of the \question command or any other lines in this document.

You can \input this .tex file right before \end{document} of your main file or compile it as a stand-alone document. Check the instructions on your conference’s website to see if you will be asked to provide this checklist with your paper or separately.

1. General Paper Structure

- 1.1. Includes a conceptual outline and/or pseudocode description of AI methods introduced (yes/partial/no/NA) **yes**
- 1.2. Clearly delineates statements that are opinions, hypothesis, and speculation from objective facts and results (yes/no) **yes**
- 1.3. Provides well-marked pedagogical references for less-familiar readers to gain background necessary to replicate the paper (yes/no) **yes**

2. Theoretical Contributions

- 2.1. Does this paper make theoretical contributions? (yes/no) **yes**, yes, please address the following points:
- 2.2. All assumptions and restrictions are stated clearly and formally (yes/partial/no) **yes**
- 2.3. All novel claims are stated formally (e.g., in theorem statements) (yes/partial/no) **yes**
- 2.4. Proofs of all novel claims are included (yes/partial/no) **yes**
- 2.5. Proof sketches or intuitions are given for complex and/or novel results (yes/partial/no) **yes**
- 2.6. Appropriate citations to theoretical tools used are given (yes/partial/no) **yes**
- 2.7. All theoretical claims are demonstrated empirically to hold (yes/partial/no/NA) **partial**

- 2.8. All experimental code used to eliminate or disprove claims is included (yes/no/NA) [yes](#)
- ### 3. Dataset Usage
- 3.1. Does this paper rely on one or more datasets? (yes/no) [yes](#), please address the following points:
- 3.2. A motivation is given for why the experiments are conducted on the selected datasets (yes/partial/no/NA) [yes](#)
- 3.3. All novel datasets introduced in this paper are included in a data appendix (yes/partial/no/NA) [NA](#)
- 3.4. All novel datasets introduced in this paper will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no/NA) [NA](#)
- 3.5. All datasets drawn from the existing literature (potentially including authors' own previously published work) are accompanied by appropriate citations (yes/no/NA) [yes](#)
- 3.6. All datasets drawn from the existing literature (potentially including authors' own previously published work) are publicly available (yes/partial/no/NA) [yes](#)
- 3.7. All datasets that are not publicly available are described in detail, with explanation why publicly available alternatives are not scientifically satisfying (yes/partial/no/NA) [NA](#)
- ### 4. Computational Experiments
- 4.1. Does this paper include computational experiments? (yes/no) [yes](#), please address the following points:
- 4.2. This paper states the number and range of values tried per (hyper-) parameter during development of the paper, along with the criterion used for selecting the final parameter setting (yes/partial/no/NA) [partial](#)
- 4.3. Any code required for pre-processing data is included in the appendix (yes/partial/no) [partial](#)
- 4.4. All source code required for conducting and analyzing the experiments is included in a code appendix (yes/partial/no) [yes](#)
- 4.5. All source code required for conducting and analyzing the experiments will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no) [yes](#)
- 4.6. All source code implementing new methods have comments detailing the implementation, with references to the paper where each step comes from (yes/partial/no) [partial](#)
- 4.7. If an algorithm depends on randomness, then the method used for setting seeds is described in a way sufficient to allow replication of results (yes/partial/no/NA) [yes](#)
- 4.8. This paper specifies the computing infrastructure used for running experiments (hardware and software), including GPU/CPU models; amount of memory; operating system; names and versions of relevant software libraries and frameworks (yes/partial/no) [partial](#)
- 4.9. This paper formally describes evaluation metrics used and explains the motivation for choosing these metrics (yes/partial/no) [yes](#)
- 4.10. This paper states the number of algorithm runs used to compute each reported result (yes/no) [yes](#)
- 4.11. Analysis of experiments goes beyond single-dimensional summaries of performance (e.g., average; median) to include measures of variation, confidence, or other distributional information (yes/no) [yes](#)
- 4.12. The significance of any improvement or decrease in performance is judged using appropriate statistical tests (e.g., Wilcoxon signed-rank) (yes/partial/no) [partial](#)
- 4.13. This paper lists all final (hyper-)parameters used for each model/algorithm in the paper's experiments (yes/partial/no/NA) [partial](#)